




(a) post-training starts here
policy cloned from 25 demos; the critic sees only these (blue), all on the band



(b) what the critic believes
learned Q ; dashed line = where it stops matching reality



(c) same critic, two updates
both follow $\nabla_a Q$; only one keeps the policy where the critic is valid



(d) measured on the red-mug task
logged residual norm during real training (log scale)

