



policy cloned from 25 demos; the critic sees only these (blue), all on the band

(a) post-training starts here




learned $\hat{Q}$; dashed line = where it stops matching reality

(b) what the critic believes



both follow $\nabla_a \hat{Q}$; only one keeps the policy where the critic is valid

(c) same critic, two updates



distance from the frozen base during led training (log scale)

(d) measured on LIBERO t65

